

Radial basis function networks

Exercise T7.1: Multi-class classification: simple methods (tutorial)

- (a) Describe how a k nearest neighbor classifier predicts the class of previously unseen inputs?
- (b) A “Parzen window” classifier extends the *electoral committee* approach of k NN. How are the different votes *weighted*?

Exercise T7.2: Radial basis function networks (tutorial)

- (a) Describe and discuss the *general architecture* of an RBF-network.
- (b) Describe and discuss the *two-step learning procedure* for RBF-networks with K basis functions.
- (c) In which cases of *regression* or *classification* do RBF networks outperform *MLPs* significantly? What are the advantages of RBF networks? In which situations are they preferable to MLP?

Exercise H7.1: Training data (homework, 1 point)

Create a sample of $p = 120$ training patterns $\{\underline{\mathbf{x}}^{(\alpha)}, y_T^{(\alpha)}\}$, $\alpha = 1, \dots, p$. The input values $\underline{\mathbf{x}}^{(\alpha)} \in \mathbb{R}^2$ should be drawn from a mixture of Gaussians with centers in an XOR-configuration according to the following scheme:

- Generate 60 samples from each of the following conditional distributions:

$$\begin{aligned} p(\underline{\mathbf{x}}|y = 0) &:= \frac{1}{2} [\mathcal{N}(\underline{\mathbf{x}}|\underline{\boldsymbol{\mu}}_1, \mathbf{I}\sigma^2) + \mathcal{N}(\underline{\mathbf{x}}|\underline{\boldsymbol{\mu}}_2, \mathbf{I}\sigma^2)] , \\ p(\underline{\mathbf{x}}|y = 1) &:= \frac{1}{2} [\mathcal{N}(\underline{\mathbf{x}}|\underline{\boldsymbol{\mu}}_3, \mathbf{I}\sigma^2) + \mathcal{N}(\underline{\mathbf{x}}|\underline{\boldsymbol{\mu}}_4, \mathbf{I}\sigma^2)] , \end{aligned}$$

with $\underline{\boldsymbol{\mu}}_1 = (0, 1)^\top$, $\underline{\boldsymbol{\mu}}_2 = (1, 0)^\top$, $\underline{\boldsymbol{\mu}}_3 = (0, 0)^\top$, $\underline{\boldsymbol{\mu}}_4 = (1, 1)^\top$

and a variance of $\sigma^2 = 0.1$.

To sample from one of the two mixture variables you can

- (i) draw with probability 1/2 whether you will draw the next point from the density in the left summand or from the density in the right summand, then
- (ii) sample from that normal distribution yielding a single point $\underline{\mathbf{x}}^{(\alpha)}$.

Note that $\mathcal{N}(\underline{\mathbf{x}}|\underline{\boldsymbol{\mu}}, \mathbf{I}\sigma^2)$ is the probability density of a multivariate normal distribution of a vector $\underline{\mathbf{x}}$, where $\underline{\boldsymbol{\mu}}$ is the mean vector and σ^2 the variance. The variance is the same for all components.

- The corresponding target values $y_T^{(\alpha)} \in \{0, 1\}$ describe the assignment to the two classes and indicate from which distribution [$p(\underline{\mathbf{x}}|y = 0)$ vs. $p(\underline{\mathbf{x}}|y = 1)$] the data point was drawn.

- (a) (1 point) Plot the resulting 120 input samples $\underline{x}^{(\alpha)}$ in a scatter plot, in which the markers and/or colors represent the corresponding samples' labels $y_T^{(\alpha)}$.

Exercise H7.2: k nearest neighbors (k NN)**(homework, 2 points)**

Build a k NN classifier that classifies new data (*query points*) by voting of the k nearest neighbors from the training set. The *electoral committee* is selected from the training patterns according to their Euclidean distance to the query point. The predicted class is determined by the target value of the majority of those k nearest patterns.

- (a) (2 points) Plot the training patterns and the decision boundary (e.g. using a contour plot or a high-resolution image of)¹ in input space for $k = 1, 3, 5$. What do you observe?

Exercise H7.3: “Parzen window” classifier**(homework, 3 points)**

This classifier implements a *weighted voting scheme*. All training points (not only the k nearest ones) cast a vote for the query point but their vote is weighted by a *Parzen window* (or *kernel function*) depending on the distance between the training samples $\underline{x}^{(\alpha)}$ and the query point \underline{x} . The Gaussian window function based on Euclidean norm $\|\cdot\|$ is:

$$\kappa(\underline{x}, \underline{x}^{(\alpha)}) = \exp\left(-\frac{1}{2\sigma_\kappa^2}\|\underline{x} - \underline{x}^{(\alpha)}\|^2\right).$$

- (a) (2 points) Plot the training patterns and the decision boundary (e.g. using a contour plot or a high-resolution image of equidistant query points) in input space for Gaussian window functions parameterized with the variances $\sigma_\kappa^2 = 0.5, 0.1$ and 0.01 .
- (b) (1 point) Add 60 new data points from a third class centered on $\tilde{\underline{\mu}} = (0.5, 0.5)^\top$ with variance $\tilde{\sigma}^2 = 0.05$. Rerun the k NN and Parzen-window classification. Plot the classification boundaries as above and compare them with your previous results.

Exercise H7.4: RBF networks**(homework, 4 points)**

Similar to the Parzen window, RBF networks classify data according to a weighted vote, but the voting committee now consists of $K \ll p$ “representatives” instead of all p data points. These representatives do not have to be previously seen data points and can be “prototypes” $\underline{t}_i \in \mathbb{R}^2$ derived from the training data via K -means clustering.

Construct an RBF network for binary classification – using the initial two-class data set from H7.1 and discarding the data points you added in H7.3b – as follows:

- Determine the K representatives \underline{t}_i via K -means clustering (you can implement the batch-algorithm described in the lecture or use an off-the-shelf implementation).
- For a given weight vector $\underline{w} \in \mathbb{R}^{K+1}$, the predicted classification for a query point \underline{x} is:

$$y(\underline{x}; \underline{w}) = \text{step}(\underline{w}^\top \underline{\phi}(\underline{x})),$$

¹Python users can generate a grid of equidistant query points that cover the range of the input variables using numpy's `meshgrid` function, then color each grid point by how the classifier's prediction.

where $\underline{\phi}(\underline{\mathbf{x}}) := \begin{pmatrix} 1 \\ \phi_1(\underline{\mathbf{x}}) \\ \vdots \\ \phi_K(\underline{\mathbf{x}}) \end{pmatrix}$ is a $(K + 1)$ -dimensional vector containing the bias and the basis function values $\phi_i(\underline{\mathbf{x}}) = \kappa(\underline{\mathbf{x}}, \underline{\mathbf{t}}_i)$ with κ from the previous exercise (Gaussian radial basis functions). Here we use the following step function

$$\text{step}(h) = \begin{cases} 1 & \text{for } h \geq 0.5, \\ 0 & \text{otherwise.} \end{cases}$$

to convert the network output $\underline{\mathbf{w}}^\top \underline{\phi}(\underline{\mathbf{x}})$ (regressed to labels 0 or 1) to a class prediction.

- Determine the weight vector as: $\underline{\mathbf{w}} = (\underline{\Phi} \underline{\Phi}^\top)^{-1} \underline{\Phi} \underline{\mathbf{y}}_{\text{True}}^\top$
where $\underline{\mathbf{y}}_{\text{True}} := (y_T^{(1)}, \dots, y_T^{(p)}) \in \mathbb{R}^{1,p}$ is the vector of target values and

$$\underline{\Phi} := (\underline{\phi}(\underline{\mathbf{x}}^{(1)}), \dots, \underline{\phi}(\underline{\mathbf{x}}^{(p)})) \in \mathbb{R}^{K+1,p}$$

- (2 points) Plot the decision boundaries together with the training patterns and locations of the representatives for $K \in \{2, 3, 4\}$. Do this for two different (reasonable²) kernel widths σ_κ of the radial basis functions ϕ_i , yielding a total of six plots.
- (2 points) We would like to visualize how the non-linearly separable data appears in the transformed feature space, which the classifier operates on:

Construct a new RBF-network with 2 RBFs and fix the centers to $\underline{\mathbf{t}}_1 = (0, 0)^\top$ and $\underline{\mathbf{t}}_2 = (1, 1)^\top$ (i.e., skip K-means clustering).

For $\sigma_\kappa = 0.45$, produce a scatter plot of the data in the space of RBF-activations, i.e. for each data point $\underline{\mathbf{x}}^{(\alpha)}$, plot $\phi_1(\underline{\mathbf{x}}^{(\alpha)})$ vs. $\phi_2(\underline{\mathbf{x}}^{(\alpha)})$, indicate their class-assignment y_T by coloring the points accordingly.

Plot also the predicted labels after training in a similar second plot.

Feel free to reduce the data-variance σ (e.g. to 0.2) to make the cluster-structure more prominent.

Total 10 points.

²You have information to make an educated guess for the value of the kernel width. The second value should be chosen to reveal a qualitative effect of a lower or larger width on the predictions.